

A Dataflow Approach to the ROS Architecture

Orestis Melkonian

Software & Knowledge Engineering Laboratory (SKEL)


o.melkonian@di.uoa.gr

December 27, 2015

Overview

- 1 Introduction
- 2 Motivation
- 3 Example
- 4 Another Example
- 5 Problems
- 6 Solutions
- 7 Related Work
- 8 Conclusion

Introduction



Future

- Ubiquitous Robotic Systems
- Web Robotics
- Minimum-cost robotics research
- Rapid prototyping


Motivation

- Satisfied with current languages on ROS?
- Right amount of abstraction?
- Logical perspective in sync with code we write?
- Real-time systems
 - Do existing languages provide time modelling?
- Continuous streams of data
 - Do existing languages provide appropriate stream handling?
- Too much technicalities that are not part of the main problem
 - Maybe a compiler can do them automatically
- Do problems faced in robotics have a dataflow nature?
 - If so, why not code in a dataflow language?

Example I

A drone progressing in a hall

- 4 sonar sensors (UP,DOWN,LEFT,RIGHT)
- Default move: AHEAD
- Depending on sensor data, change direction to avoid collision with any of the four walls
- Also occasional sonar sensor malfunction (negative values)



Example II

```
float arr[4] = {0, 0, 0, 0};
void callback_up(float val) { if (val > 0) arr[0] = val; }
void callback_down(float val) { if (val > 0) arr[1] = val; }
void callback_left(float val) { if (val > 0) arr[2] = val; }
void callback_right(float val) {
    if (val > 0) arr[3] = val;
    string cmd_ver = "", cmd_hor = "";

    if (arr[0] > arr[1]) cmd_ver = "UP"
    else if (arr[0] < arr[1]) cmd_ver = "DOWN"
    else cmd_ver = ""
    if (arr[2] > arr[3]) cmd_hor = "LEFT"
    else if (arr[2] < arr[3]) cmd_hor = "RIGHT"
    else cmd_hor = ""

    ROS.publish("cmd", cmd_ver + cmd_hor);
}


int main() {
    ROS.subscribe("sensor_up", callback_up); ROS.subscribe("sensor_down", callback_down);
    ROS.subscribe("sensor_left", callback_left); ROS.subscribe("sensor_right", callback_right);
    ROS.publisher("cmd");
    ...
}
```

Example III

```
Topic fused = zip
    (subscribe "sensor_up",
     subscribe "sensor_down",
     subscribe "sensor_left",
     subscribe "sensor_right")
    >>> filter (> 0)

Node actor = publish "cmd" (map act fused)

act (u,d,l,r) = ver ++ hor where
    ver = if u>d "UP" else (if u<d "DOWN" else "")
    hor = if l>r "LEFT" else (if l<r "RIGHT" else "")
```




Another Example

Implementation of a PID controller

```

constant Kp, Ki, Kd, setpoint
Graph g = loop(output):
    error = setpoint - output
    p = (*Kp) <<< error
    i = (*Ki) <<< integral <<< error
    d = (*Kd) <<< derivative <<< error
    result = sum [p, i, d]
    output <<< process <<< result
  
```



Problems I

■ Scalability

- Asynchronous behaviour using callback functions
- Complex schemes require "internal plumbing"
- Unreadable, difficult to maintain
- Does not separate data from control

■ Untyped topics

- Topics are just named pipes
- No type-safety
- Hard to impose constraints on them
- e.g. Visual detection on audio streams
 - ▶ Does not catch errors at compile time

Problems II

■ Time modelling

- Control theory models robot behaviour primarily using differential equations on time
- Existing languages do not provide the concept of time
- Need to implement tedious delta timing
- But why? Time is logical
- Need to have temporal nature by design

■ Compiler restriction


- Coder explicitly specifies where nodes are being executed
- Won't allow intelligent runtime re-configurations
e.g. for optimal power management, minimum execution time

Problems III

■ Verifiability


- Robots are becoming an integral part of our everyday life
- More and more responsibility
e.g. babysitters, security monitoring, space robots
- Must be able to prove program correctness
- Nearly impossible for low-level languages
- Need for more precise semantics

Solutions I




■ A domain-specific dataflow language

- ROS architecture essentially defines a dataflow graph



Solutions II

- Composable stream-processing operators



Solutions III


- Timely model \rightarrow Simpler, portable code

e.g.

$$x(t) = 1/2 \int_0^t (u_l(t) + u_r(t)) \cos \theta(t) dt$$

\Downarrow

```
xs = (v &&& t) >>> * >>> integral >>> /2
where v = (vr &&& vl) >>> +
      t = theta >>> cos
```



Solutions IV

- Separates data from control

```
void handle_data(float val) {  
    if (val > threshold)  
        publish("cmd", act(val));  
}  
  
int main() {  
    ROS.subscribe("sensor", handle_data);  
    ROS.publish("cmd");  
    ...  
}
```



```
Topic myTopic = subscribe "sensor" >>> filter(> threshold)  
Node actor = publish "cmd" (map act myTopic)
```


Solutions V

■ Topics as streams

- First-class citizens
- At last, type-safety!

```
Topic t :: Stream[Float]
```

```
Topic t = subscribe "sonar1" >>> filter(> 0)
```

- Automatic *rosmmsg* generation
- We now can catch errors at compile time

■ Pure functions to the rescue

- Keep maximum portion of the program as pure functions
- No side-effects → equational reasoning
- Can prove program correctness

Solutions VI

■ **Dynamic reconfiguration**

- Locality-agnostic node execution
- System will allocate nodes to machines for optimality
- Ability to specify preferences
e.g. `run 'graph' --powersave/--uniform /--mintime`

■ **Harness the cloud!**

- Adopt the advantages of web-based solutions to develop robotic applications
- Use the cloud to get access to distributed computing resources
e.g. Grasping using the Google Object Recognition Engine

Functional Reactive Programming (FRP)

- A programming paradigm suitable for hybrid systems
- Uses the building blocks of functional programming (*map, filter, ...*)
- For systems that are:
 - Responsive
 - Resilient
 - Elastic
 - Message-driven
- Explicit time modelling
- Express event-handling in a natural way
- Applications in GUIs, robotics and music.

Robotics

■ roshask

- A Haskell-binding client library for ROS
- Takes the first step towards stream-oriented robotics
- Implements topics as first-class citizens
- Provides some basic stream manipulation

■ Frob

- A domain-specific language (DSL)
- Realization of the FRP model on robotics
- Complex robot behaviours in a few lines of code
- Implementation exists in Haskell (Yampa)
- Several performance issues (time-/space-leaks)

Big Data

■ Akka

- A framework for scalable, distributed, message-driven applications
- Incorporates the Actor model, like ROS
- The stream handling seemed to be error-prone and tedious
- The inability to treat streams efficiently and intuitively led them to develop the Akka Streams API
- Also influenced by FRP

■ Naiad

- A distributed system for cyclic parallel dataflow programming
- A new computational model: *Timely dataflow*
- Many novel approaches for efficient dataflow graphs

Ziria: A DSL for wireless systems programming

- Designed by Haskell-lovers
- Also inspired by Functional Reactive Programming (FRP)
- Towards Software-Defined Radios (SDR)
- In contrast to SORA (low-level C++ library)
- High-level functional language replaces low-level C
- Implementation of a Wifi Receiver
 - SORA → 23000 lines of code
 - Ziria → 3000 lines of code
- Implementation of a Scrambler
 - SORA → 90 lines of code
 - Ziria → 20 lines of code
- But same performance!

Conclusion

- There will be some learning curve
- But it will pay off in terms of:
 - Productivity
 - Cleaner, safer, scalable code
 - More efficient resource management
 - Faster robotics research
- Any DSL design relies heavily on feedback from domain experts
- There is a general pattern here
 - Abstraction is the way!

The End