

THE EXTENDED UTXO MODEL

Manuel M.T. Chakravarty, James Chapman, Kenneth MacKenzie,
Orestis Melkonian, Michael Peyton Jones, Philip Wadler
(presented by **Alexander Nemish**)

February 14, 2020



INTRODUCTION

BACKGROUND

Bitcoin



- Based on *unspent transaction outputs* (UTXO)
- Simple, non Turing-complete smart contracts in SCRIPT

Ethereum



- Account-based
- Turing-complete smart contracts in Solidity/EVM

Cardano (IOHK)



- UTXO-based + extensions
- Turing-complete smart contracts in Plutus



- Focus on validating the relevant meta-theory
 - In contrast to validating individual contracts
- Fully mechanized approach, utilizing Agda's rich type system
- Fits well with IOHK's research-oriented approach



- Detailed description of the Extended UTXO model (EUTXO)
- Formalization in Agda
- Formal proof of bisimulation with a specific form of state machines

EUTXO, INFORMALLY...

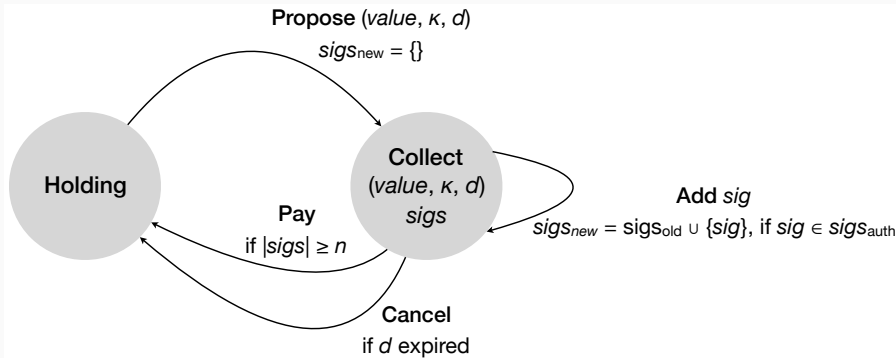
UTXO vs EUTXO



- New data value on outputs
- More information available to validators
- We show that Cardano can model a specific form of state machine
 - However, much more computational patterns are possible
 - e.g. the entirety of **Marlowe**, a DSL for financial contracts, has been implemented as a state machine on top of EUTXO.

EXAMPLE: MULTI-SIGNATURE CONTRACT

- **n-out-of-m** signature scheme
- Plain UTXO requires off-chain communication
- Can be expressed as a simple state machine:



EXAMPLE: IMPLEMENTATION IN EUTXO

- State machine is associated to a *validator* function
- *Data values* in outputs correspond to states
 - $\in \{\text{Holding, Collecting}\}$
- Validator takes one of four possible transitions
 - $\in \{\text{Propose, Add, Cancel, Pay}\}$
 - Choice provided by the *redeemer* of the spending input

EUTXO, FORMALLY...

1. **Data values** additionally carried by outputs
 - passed as extra argument of type `Data` during validation
 - allows a contract to carry data without changing its code
 - otherwise we would not be able to identify a contract by its code's hash
2. More information about the transaction available to the validator
 - passed as extra argument of type `TxInfo` during validation
 - allows inspection of the transaction's outputs, thus supporting **contract continuity** (i.e. outputs target the expected validator)

3. Transactions have (restricted) access to time

- addition transaction field: **validity interval**
- specifies a time interval, in which the transaction must be processed
- in contrast to allowing access to the current time
 - allows for **deterministic** script execution
 - we can pre-calculate consumed resources/time
 - a user can simulate execution locally

Quantity	an amount of currency
Tick	a tick
Address	an “address” in the blockchain
Data	a type of structured data
DataHash	the hash of a value of type Data
TxId	the identifier of a transaction
txId : Tx → TxId	get a transaction’s identifier
Script	the (opaque) type of scripts
scriptAddr : Script → Address	the address of a script (i.e. its hash)
dataHash : Data → DataHash	the hash of a data value
[[_]] : Script → Args → Bool	applying a script to its arguments

DEFINED TYPES

Output = (*value* : Quantity, *addr* : Address, *dataHash* : DataHash)

OutputRef = (*id* : TxId, *index* : \mathbb{N})

Input = (*outputRef* : OutputRef, *validator* : Script,
data : Data, *redeemer* : Data)

Tx = (*inputs* : Set[Input], *outputs* : List[Output],
validityInterval : Interval[Tick])

Ledger = List[Tx]

VALIDITY OF EUTXO TRANSACTIONS (I)

1. The current tick is within the validity interval

$$\text{currentTick} \in t.\text{validityInterval}$$

2. All outputs have non-negative values

$$\forall o \in t.\text{outputs}, o.\text{value} \geq 0$$

3. All inputs refer to unspent outputs

$$t.\text{inputs} \subseteq \text{unspentOutputs}(l)$$

4. Value is preserved (ignoring fees)

$$\sum_{i \in t.\text{inputs}} \text{getSpentOutput}(i, l).\text{value} = \sum_{o \in t.\text{outputs}} o.\text{value}$$

5. No output is double spent

$$\text{If } i_1, i_2 \in t.\text{inputs} \text{ and } i_1.\text{outputRef} = i_2.\text{outputRef} \text{ then } i_1 = i_2$$

6. All inputs validate

$$\forall i \in t.inputs, \llbracket i.validator \rrbracket(i.data, i.redeemer, toTxInfo(t, i)) = \text{true}$$

7. Validator scripts match output addresses

$$\forall i \in t.inputs, \text{scriptAddr}(i.validator) = \text{getSpentOutput}(i, l).addr$$

8. Data values match output hashes

$$\forall i \in t.inputs, \text{dataHash}(i.data) = \text{getSpentOutput}(i, l).dataHash$$

EXPRESSIVENESS OF EUTXO

CONSTRAINT EMITTING MACHINES (CEM)

To formally reason about the expressiveness of EUTXO, we introduce a specific form of state machines:

- Type of states S , type of inputs I
- $\text{final} : S \rightarrow \text{Bool}$
- $\text{step} : S \rightarrow I \rightarrow \text{Maybe } (S \times \text{TxConstraints})$

These are similar to Mealy machines, but differ in some aspects:

1. No notion of initial states
2. Cannot transition out of a final state
3. Blockchain-specific output values (TxConstraints)
 - these typically are first-order equality constraints on the fields of Tx

- A ledger l corresponds to a CEM state s :

$$l \sim s$$

- New (valid) transaction submitted to ledger l :

$$l \xrightarrow{tx} l'$$

- Valid CEM transition from source state s to target state s' , using input symbol i and emitting constraints tx^{\equiv} :

$$s \xrightarrow{i} (s', tx^{\equiv})$$

Given a smart contract, expressed as a CEM C , we can derive the validator script that disallows any invalid transitions:

$$\text{validator}_C(s, i, txInfo) = \begin{cases} \text{true} & \text{if } s \xrightarrow{i} (s', tx^{\equiv}) \\ & \text{and satisfies}(txInfo, tx^{\equiv}) \\ & \text{and checkOutputs}(s', txInfo) \\ \text{false} & \text{otherwise} \end{cases}$$

BEHAVIOURAL EQUIVALENCE: WEAK BISIMULATION

Proposition 1 (Soundness)

Given a valid CEM transition, we can construct a new valid transaction, such that the resulting ledger corresponds to the target CEM state:

$$\frac{s \xrightarrow{i} (s', tx^{\equiv}) \quad l \sim s}{\exists tx \, l' . l \xrightarrow{tx} l' \wedge l' \sim s'} \text{ SOUND}$$

Proposition 2 (Completeness)

Given a new valid transaction on the ledger, it is either irrelevant to the state machine or corresponds to a valid CEM transition:

$$\frac{l \xrightarrow{tx} l' \quad l \sim s}{l' \sim s \vee \exists i \, s' \, tx^{\equiv} . s \xrightarrow{i} (s', tx^{\equiv})} \text{ COMPLETE}$$

RELATED WORK

- **Bitcoin Covenants**
 - Allows restricting how output values will be used in the future
 - Major inspiration for our introduction of *data values*
- **Bitcoin Modelling Language (BitML)**
 - A process-calculus for Bitcoin smart contracts, whose operational semantics comprise a state machine
 - Compiles down to Bitcoin transactions, without any extensions
 - Quite complicated translation and requires off-chain communication
- **Scilla**
 - For Ethereum contracts, using message-passing state machines dubbed *Communicating State Transition Systems*
 - Embedded in Coq, hence amendable to formal verification
 - Allows proving temporal (hyper-)properties
- **Bitcoin Contracts as Timed Automata**
 - Pragmatic model checking using UPPAAL
 - Does not come with formal guarantees though

QUESTIONS?